

# Debugging web applications in a production environment


strace, ps and gdb.

# Overview

- Mindset
- Production Environments
- Web applications
- Workflow
- Strace, ps and gdb
- Lessons learned
- Last production debugging session:
  - Lasted a number of days
  - Peaked at ~400 orders/second across ~90 servers
  - *Gigabytes* of data collected.

**Mindset**

'It is a capital mistake to theorize before one has data. Insensibly one begins to twist facts to suit theories, instead of theories to suit facts.'




**'Eliminate all other factors, and the one which remains must be the truth.'**




# Production Environments

# frogger



- Users don't stop
- Can't edit the code, no print/warn
- Can't see what they see (especially in SSL environment)
- Security concerns
- Fast!
- Data overload!

# Work flow





# Nature of web apps

Dr. Watson, it's a well known fact that 87.36% of web applications are I/O bound

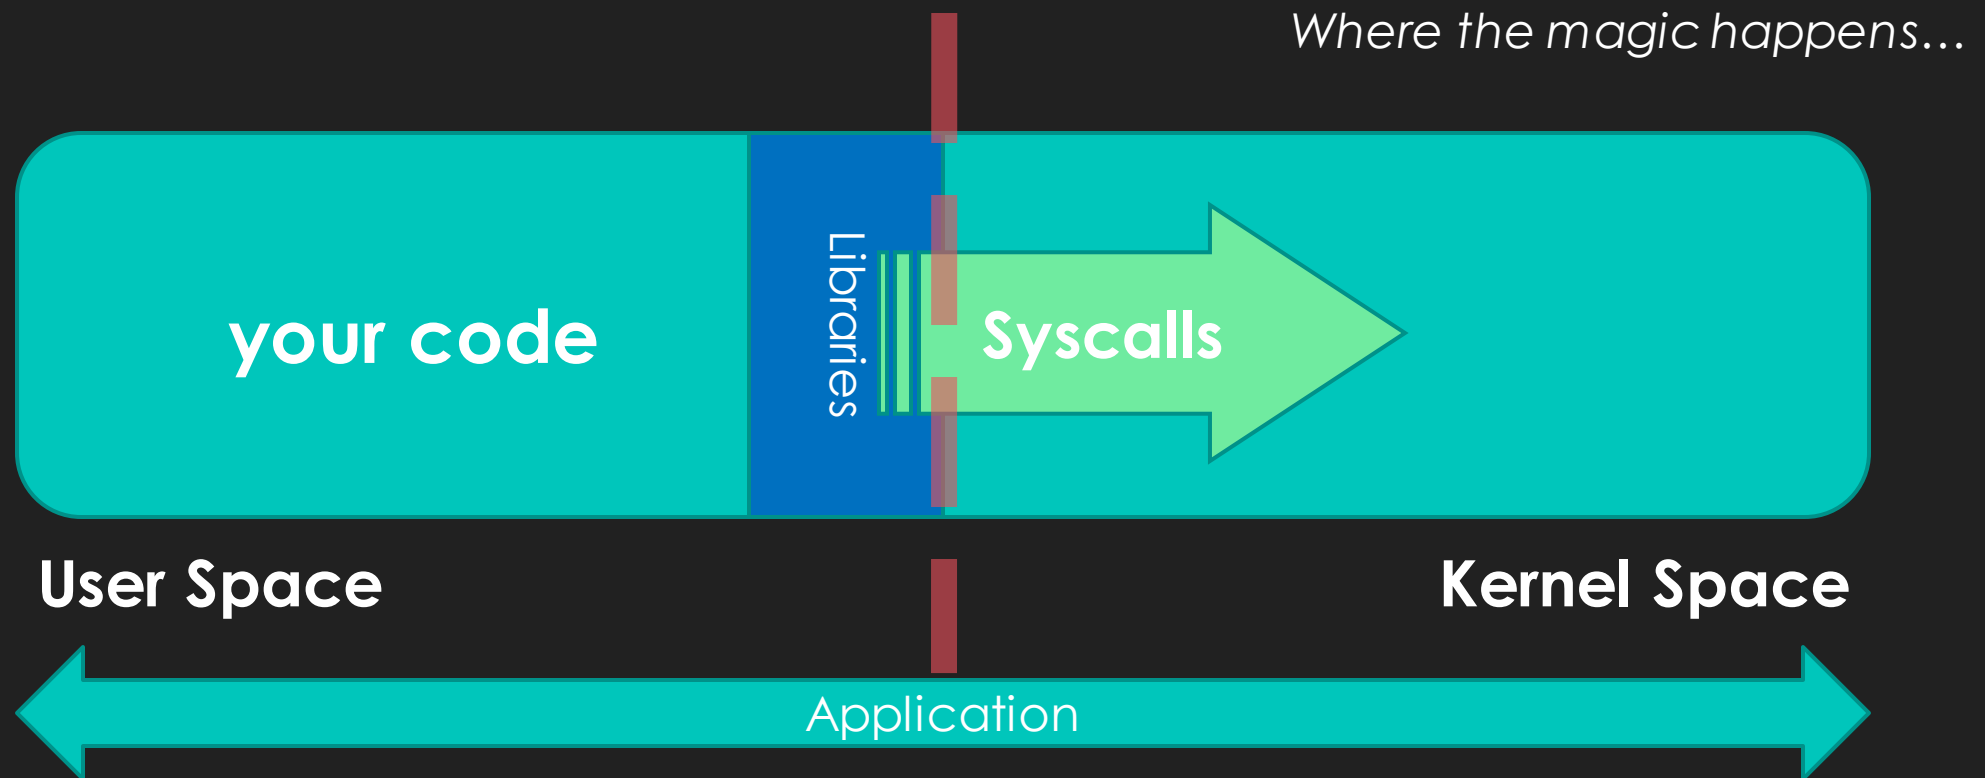
Sherlock Holmes



# Making sausage.

- You take stuff in various sources
- You output a web page
- System calls are your I/O interface.

# System calls



# System calls

## Common System calls

- `open()`
- `read()`
- `write()`
- `execve()`
- `brk()`
- `time()`
- `mmap()`
- `ioctl()`
- `connect()`

## More Info:

- System Calls are described in section 2 of 'man'
  - E.g `man 2 mmap2`
- System calls are C functions and will be of the format `function(arg1, arg2, .. argN) → value`

**Strace**

# strace intro

- Allow you to see which syscalls an application is executing along with arguments and return values
- Two Modes
  1. Run the application with strace
  2. Attach to a running application (with `-p`.) **Needs root access!**
- Can be found in your Linux distribution's package manager (`apt-get`, `yum`, etc.)
- Generates a LOT of data.
  - 116 lines of output for a simple `"ls"`
  - Filter, filter, filter, `grep`.

**strace sample**



# Strace options

- -c : show summary table
- -f : follow child processes
- -T : show time spent in system calls
- -t, -tt, -ttt : show start time, with millis or in epoch time
- -e trace=[set name] : trace only a certain set of calls.
- -e read=fd1, fd2, etc write=fd1,f2 : show ascii and hex dump of data read/written
- In order to use ssq.pl please use: -Ttt -f
- More options and examples listed in the man page

# What to look for

- Long running system calls
  - Slow network activity
  - Bad hardware
  - Bad Cable!
- Infinite/Long Running Loops (repeated syscalls doing the same thing.)
- Unnecessary I/O:
  - Possible cache candidates
  - Reading large, unexpected datasets (e.g. an SQL query gone wrong.)
- Customer's view.

# Ssql.pl

- Download from <https://github.com/jonphilpott/stracetools>
- Very helpful for Filter and Analyze stages:

```
sqlite> select * from strace order by dur desc limit 5 ;
```

```
1177|17061|22:57:45|136849|read|read(0, "l", 1)|1|2.165879
```

```
1398|17061|22:57:48|172280|read|read(0, "\r", 1)|1|0.910384
```

```
1474|17061|22:57:49|96344|read|read(0, "\4", 1)|1|0.666353
```

```
1197|17061|22:57:47|704137|read|read(0, "\r", 1)|1|0.428623
```

```
1192|17061|22:57:47|473728|read|read(0, "&", 1)|1|0.229005
```

ps

# Regular ps output

```
jon@jon-craptop:~/stracetools$ ps -eal|head
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
4	S	0	1	0	0	80	0	-	917	poll_s	?	00:00:01	init
1	S	0	2	0	0	80	0	-	0	kthrea	?	00:00:00	kthreadd
1	S	0	3	2	0	80	0	-	0	run_ks	?	00:00:35	ksoftirqd/0
1	S	0	6	2	0	-40	-	-	0	cpu_st	?	00:00:00	migration/0
1	S	0	7	2	0	-40	-	-	0	watchd	?	00:00:07	watchdog/0
1	S	0	8	2	0	-40	-	-	0	cpu_st	?	00:00:00	migration/1
1	S	0	10	2	0	80	0	-	0	run_ks	?	00:00:13	ksoftirqd/1
1	S	0	11	2	0	-40	-	-	0	watchd	?	00:00:06	watchdog/1
1	S	0	12	2	0	-40	-	-	0	cpu_st	?	00:00:00	migration/2



Process state

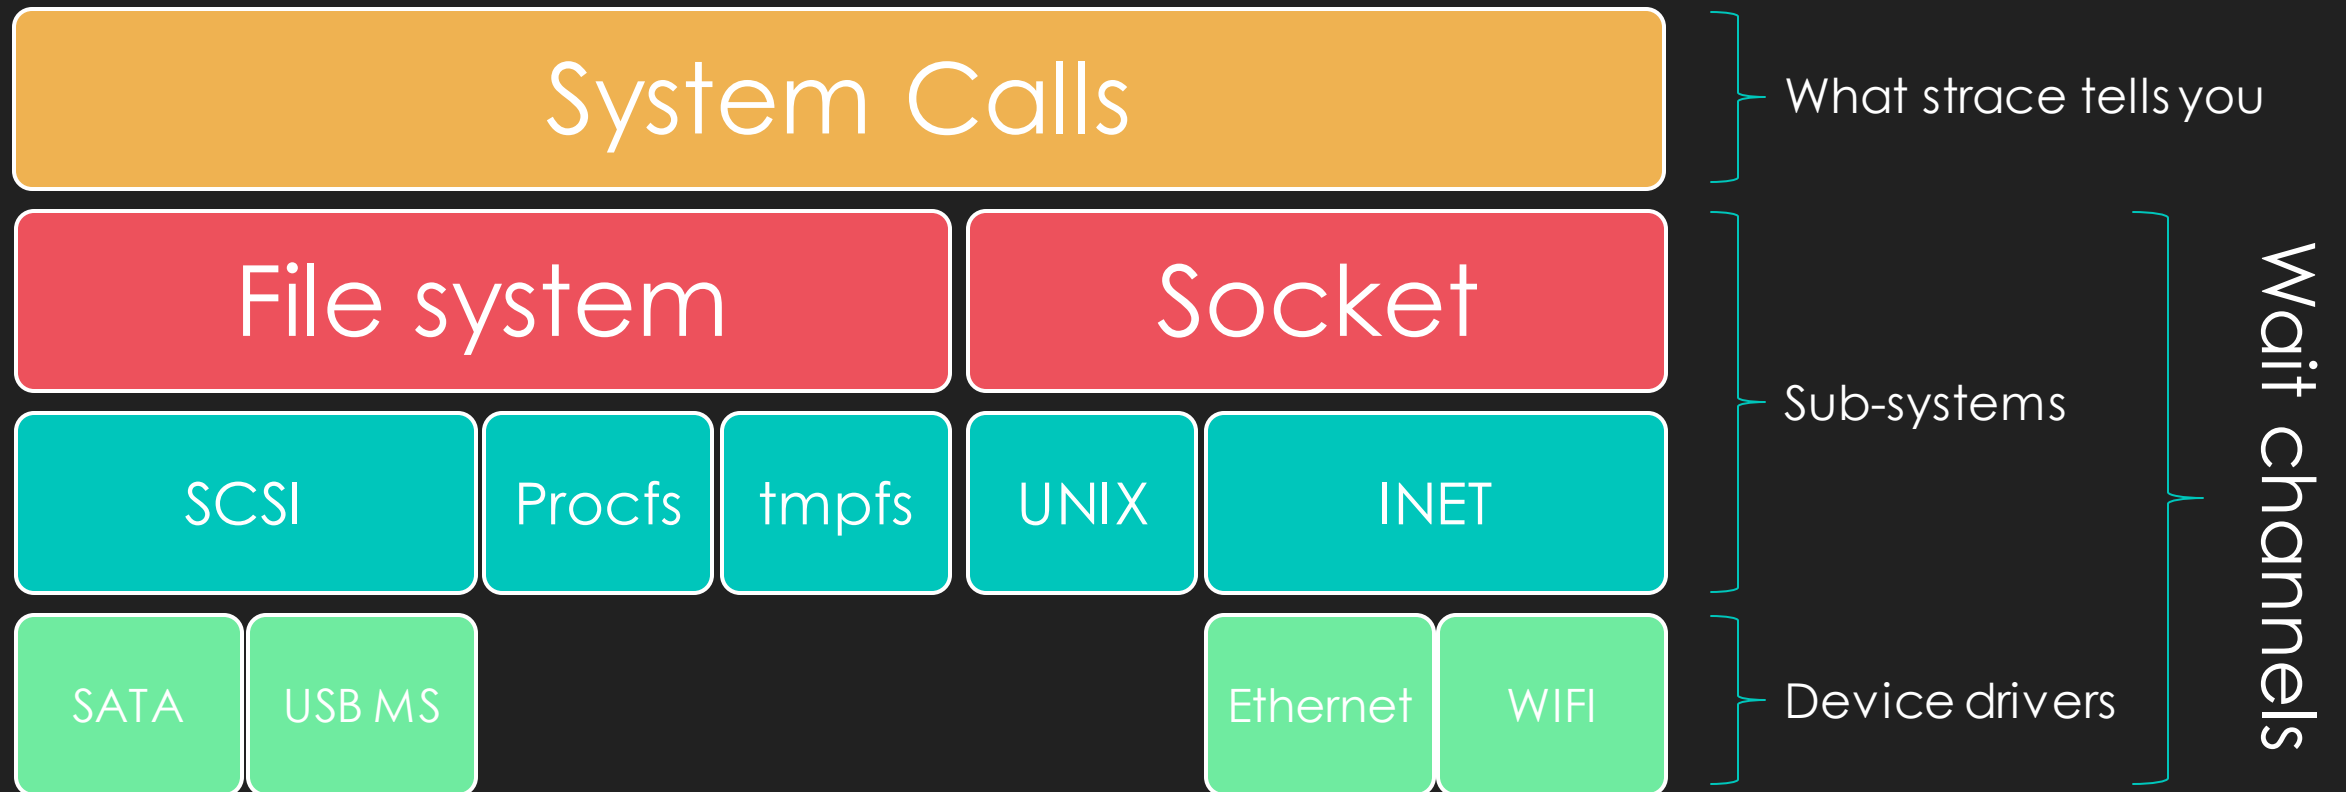


Wait channel

# Process states

STATE	DESCRIPTION
<b>D</b>	<b>UNINTERRUPTIBLE SLEEP (I/O) –ARGH!!!!</b>
<b>R</b>	Running/Runnable
<b>S</b>	Interruptible sleep
<b>T</b>	Stopped
<b>W</b>	Paging (not in 2.6 kernels)
<b>X</b>	Dead (never see this)
<b>Z</b>	Zombie

# Wait Channels



# wtf Channels?

- Tell you where *exactly* your process is stuck in kernel land.
- Two places to find out what the wait channel means
  1. The sensible name of the function, e.g.:  
`scsi_error_handler`
  2. The kernel source and grep.



# procmon.pl

- Download from <https://github.com/jonphilpott/stracetools>
- Allows you collect process information over time
- Can output CSV or into a Sqlite3 database.

```
sqlite> select * from ps where vmswap > 0 limit 5;  
3|1|1365654779|init|S|0|0|3668|3672|1892|712|136|184|2532|104|1|poll_schedule_timeout  
51|799|1365654779|upstart|S|0|0|2848|2848|288|176|136|132|2332|56|1|poll_schedule_timeout  
54|908|1365654779|dbus|S|102|105|6172|6212|3560|3096|136|424|2404|172|1|poll_schedule_timeout  
55|920|1365654779|modem|S|0|0|7224|7288|2076|400|136|388|5988|368|1|poll_schedule_timeout  
56|925|1365654779|bluetoothd|S|0|0|4744|4744|1196|180|136|840|3456|156|1|poll_schedule_timeout
```

**gdb**

The GNU debugger is  
the only tool I use.

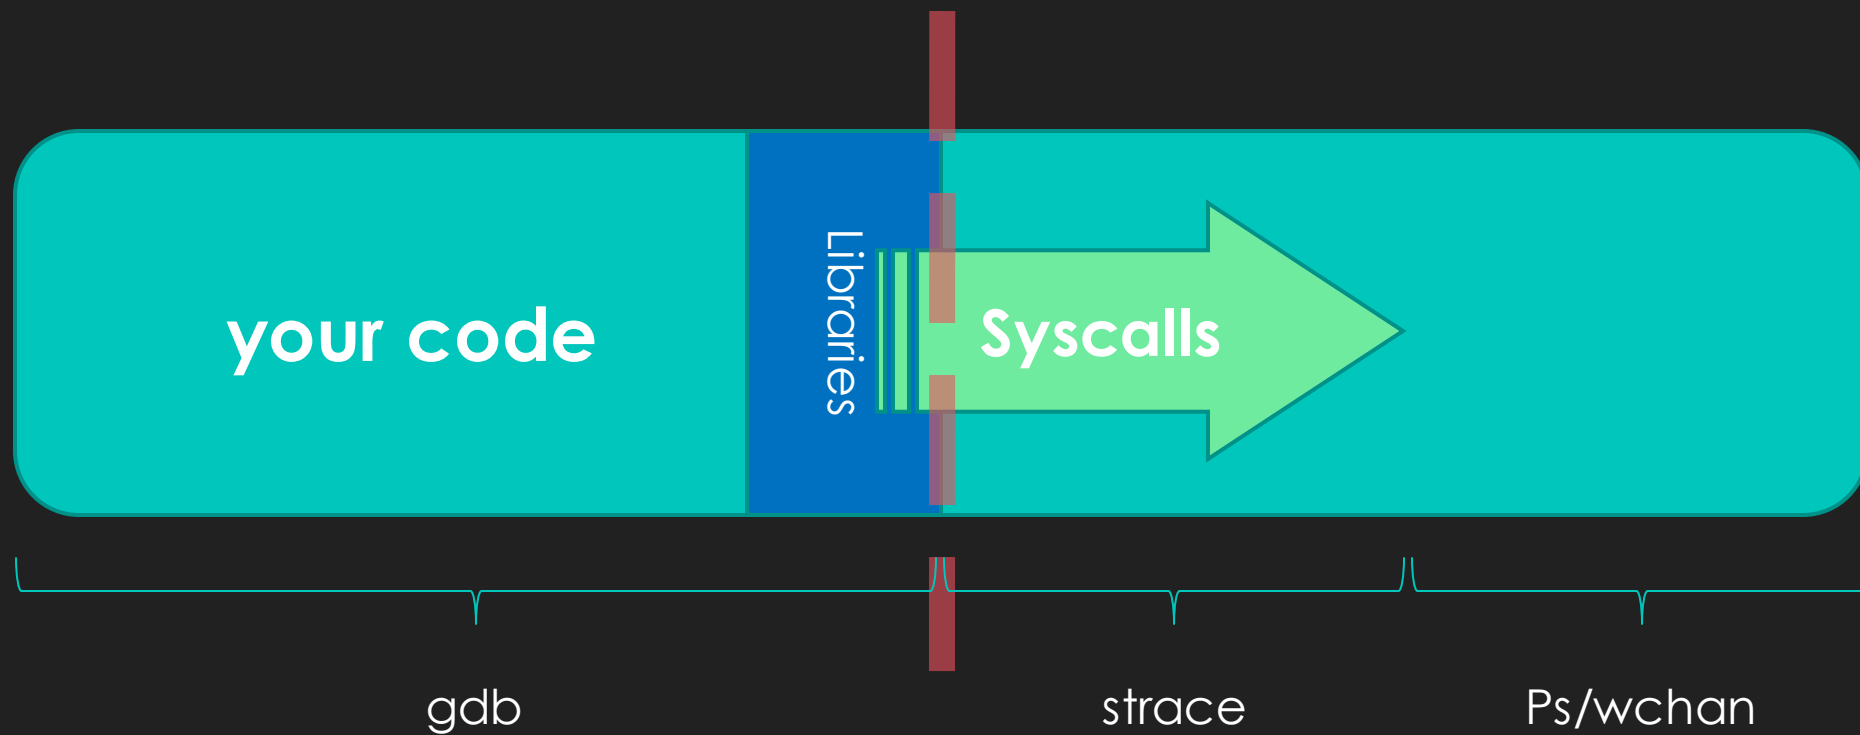
Sherlock Holmes



# gdb – The GNU debugger

- Like strace; has two modes
  - Execute the application within gdb
  - Attach to an existing process
- The ultimate tool
  - And the hardest to use
- The last resort:
  - Pauses processes

# Where gdb sits



# gdb simple tools

- Simple techniques to gather information
  1. Controlling execution
  2. The Stack
  3. Examining and altering data
  4. Breakpoints
  5. Macros, logging and .gdbinit
  6. Executing code.

C pointer syntax?  
Oh hell, no.

Dr. Watson



# gdb commands

- `continue` : Continue Execution
- `finish` : execute until the end of the current stack frame
- `backtrace` : display stack trace
- `up/down #` : move up and down the stack
- `x/16wx $esp` : Display 16 words of data at the top of the stack
- `x/s *(char **)( $esp + n*4)` : display nth string argument
- `Info registers`
- `*(int **)( $esp + n * 4) = 1234;` : set value to 1234
- `set $ret_value = function(args, ...)` : execute C function return value into \$ret\_value
- `break function` : set breakpoint on function
- `catch syscall systemcallname`
- `continue` : continue execution
- `Finish` : run current functions until completion
- `Condition n [conditional]` : set breakpoint condition



# **gdb macros**

```
define xna
```

```
    x/s *(char **)($esp + $arg0 * 4)
```

```
end
```

# Executing perl code

- Ancient Perl (5.8.4 + No Symbols!)

```
define run_perl
    set $sv = Perl_eval_pv($arg0, 1)
    printf "ret: %s\n", **$sv
end
```

- Modern Perl (Needs Symbols ☹)

```
define run_perl
    set $sv =
Perl_eval_pv(PL_curinterp, $arg0, 1)
    printf "%s", *(char **)((void
*)($sv + 12))
end
```

# Play with it



- Take some time and tinker
- Dig
- Break it
- ... but not in prod.

# Questions?

- Email: [jonathan.philpott@ticketmaster.com](mailto:jonathan.philpott@ticketmaster.com)
- Presentation and tools downloadable from:  
<https://github.com/jonphilpott/stracetools>